



Artificial Intelligence and Data Analytics for Engineers (AIDAE)

Lecture 8
June, 26th

Andrés Posada
Anas Abdelrazeq
Marco Kemmerling
Vladimir Samsonov

Today's Lecturer

Artificial Intelligence and Data Analytics for Engineers

Overview Lectures 1 – 4



1

Introduction to Data Analytics and Artificial Intelligence in Engineering: Organizational matters (e.g. exam, exercises, dates). Goals, Challenges, Obstacles, and Processes.



2

Introduction into the primary programming language of the lecture, Python: Syntax, libraries, IDEs etc. Why is Python the *lingua franca* of the Data Scientist?



3

Data Preparation: Cleansing and Transformation. How do real world data sets look like and why is cleaning and transformation an integral part of a Data Scientist's workflow?



4

Data Integration: Architectures, Challenges, and Approaches. How can you integrate various data sources into an overarching consolidating schema and why is this important?



5

Data Representation: Feature Extraction and Selection. How to pick relevant features for the task at hand. Manual vs automatic methods. What is the curse of dimensionality?



6

Data-Driven Learning: Supervised (Classification, Regression) methods and algorithms. What is an artificial neural net? What methods are there for evaluation of your model?



7

Data-Driven Learning: Unsupervised (Clustering) methods and algorithms. How can machines learn without labels? What methods are there for evaluation of your model?

8

Environment-Driven Learning: Reinforcement Learning

Today's Lecture

Reinforcement Learning

What is it?

Methods

Applying

Learning Objectives



Learning Objective
w.r.t. Knowledge/Understanding.

After successfully completing this lecture, the students will have achieved the following learning outcomes:

- Have an understanding of what reinforcement learning is.
- Understand topics like agents, rewards, action.
- Know about the different types of reinforcement learning.

Motivation and Introduction

*„Artificial Intelligence =
Deep Learning + Reinforcement Learning“,
David Silver, Google DeepMind 2015*



- **Deep Learning:** Learning an abstract representation of data
- **Reinforcement Learning:** Learning parameters of representation from interaction with the environment


Reinforcement Learning in Action

Adaptive Path Planning




Reinforcement Learning in Action

Human-Level Dexterity with Robotics



Unperturbed



With perturbations

Solving complex manipulation and planning tasks in unknown conditions while been trained in the simulation only

source: <https://openai.com/blog/solving-rubiks-cube/>

Introduction

Introduction

- Brief view into Artificial intelligence and Machine Learning

Artificial Intelligence:
Any technique which enables a computer to mimic human behaviour.

General AI (GAI)
Transfer knowledge
across domains.

Narrow AI
Performs a single task extremely well

Machine Learning
Algorithms whose performance improve as they are exposed to more data.

Supervised Learning

Unsupervised
Learning

Reinforced Learning


What is reinforcement learning?



Main Idea

The agent learns from success and failure through reward and punishment.

- Goal of the agent is defined by the reward function.
- Agent receives Feedback in form of a reward
With RL, the agent does not know the reward function!
- Agent must (learn) to act in such a way that the expected future reward is maximized
- All learning is based on samples of states, actions and rewards




Essential Aspects of Reinforcement Learning



Reinforcement learning is formulated as a Markov decision-making process (MDP). This process is a tuple $\langle S, A, P, R, \gamma \rangle$

Example:



- S is a finite set of states
- A is a finite set of actions
- P is a transition matrix
 - $P_{ss'}^a = \mathbb{P}[S_{t+1} = s' | S_t = s, A_t = a]$
 - Probability that the action a leads from the state s to the state s'
- R is a reward function, $R_s^a = \mathbb{E}[R_{t+1} | S_t = s, A_t = a]$
- γ is a discount factor $\gamma \in [0, 1]$



Definition:

Markov Property:

$$P[S_{t+1} | S_t] = P[S_{t+1} | S_1, \dots, S_t]$$

Where:

S_t – current state of the agent

S_{t+1} – next state

P – state transition probability

In other words:

Future is Independent of the past given the present



Definition:

Expected Return – is a sum of future returns:

$$G_t = R_{t+1} + \gamma G_{t+1}$$

Where:

G_t – expected return of the current state

R_{t+1} – reward of the next state

γ – discount factor

How to understand this equation:

$$G_1 = R_2 + \gamma R_3 \dots + \gamma^8 R_{10} = R_2 + \gamma(R_3 \dots + \gamma^7 R_{10})$$

$$G_2 = R_3 \dots + \gamma^7 R_{10}$$

$$G_1 = R_2 + \gamma G_2$$



Definition:

Value Function (State Value Function) – expected return for policy π in state S at time t

$$V_{\pi}(S_t) = E_{\pi}(G_t | S_t)$$

Where:

S_t – current state

π – policy RL agent follows

G_t – expected return for the current state

How to understand E_{π} :

- Actions taken by the agent can be stochastic
- The transition probabilities in real world environment can be stochastic

We need to account for probability of every possible transition in the current state:

$$E_{\pi}(G_t | S_t) = \sum_{a \in A} P(a | S_t) \sum_{S_{t+1} \in S} \sum_{R_{t+1} \in R} P(S_{t+1}, R_{t+1} | S_t, a) G_t$$

Where: S_t, S_{t+1} - current and next states; a - agent action, A, S, R - action, state and reward spaces; P denote conditional probabilities



Definition:

Q-Function (Action Value Function) – expected return for policy π in state S at time t assuming action a

$$Q_{\pi}(S_t, a) = E_{\pi}(G_t | S_t, a)$$

Where:

S_t – current state

π – policy RL agent follows

a – agent action

G_t – expected return for the current state



Definition:

Bellman Equation – calculating the value of the current state depends only on the possible next states

$$V_{\pi}(S_t) = E_{\pi} (R_{t+1} + \gamma V_{\pi}(S_{t+1}) | S_t)$$

Where:

S_t – current state

S_{t+1} – next state

R_{t+1} – reward of the next state

γ – discount factor

π – policy RL agent follows

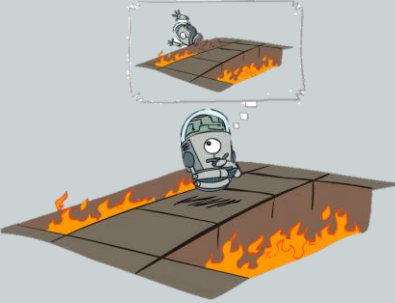

If environment transitions and RL policy are deterministic:

$$V_{\pi}(S_t) = R_{t+1} + \gamma V_{\pi}(S_{t+1}) | S_t$$

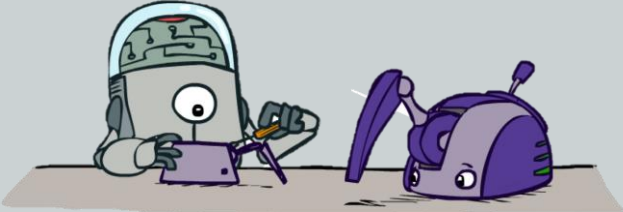
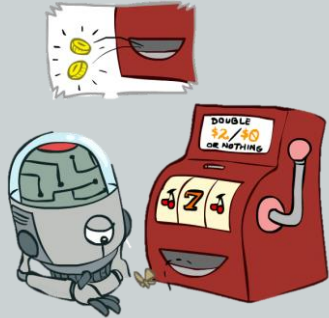
Most of the model-free RL algorithms are just trying to solve Bellman Equation!

Types of RL

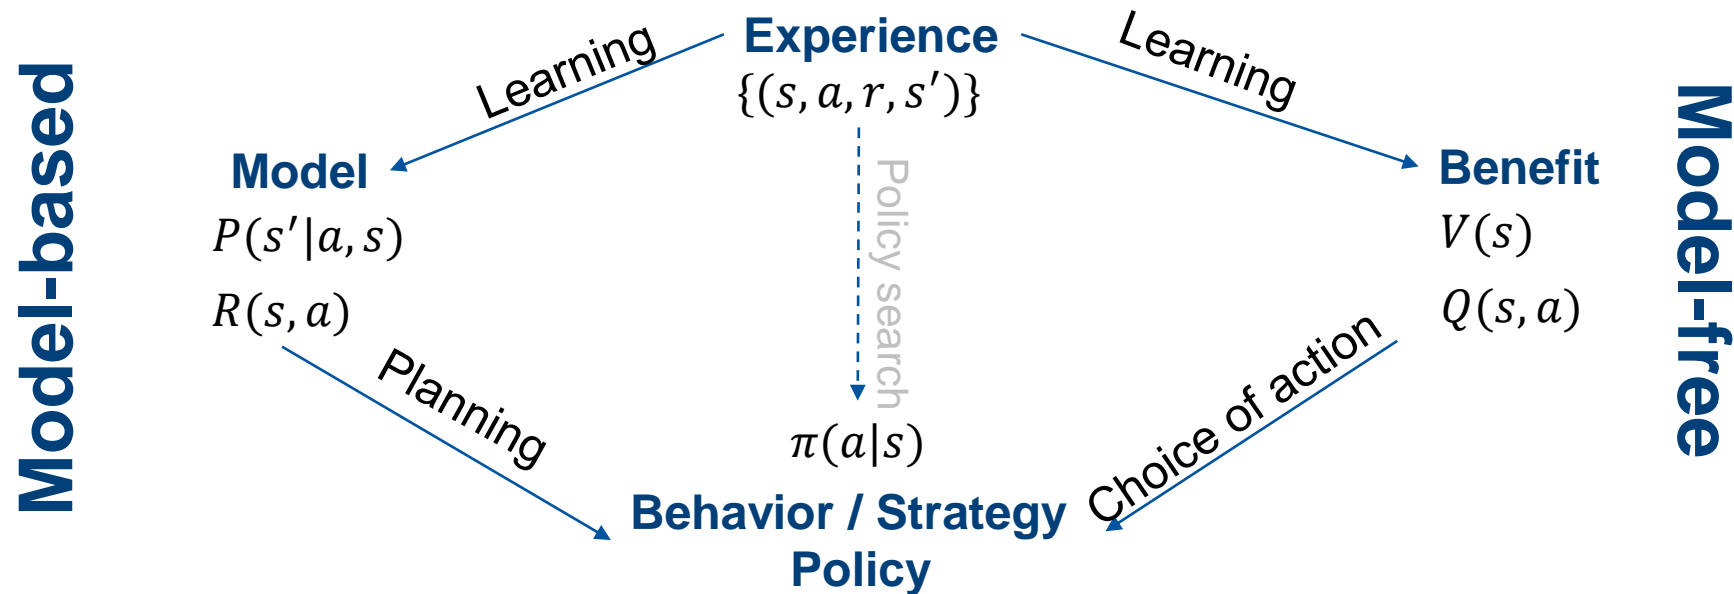
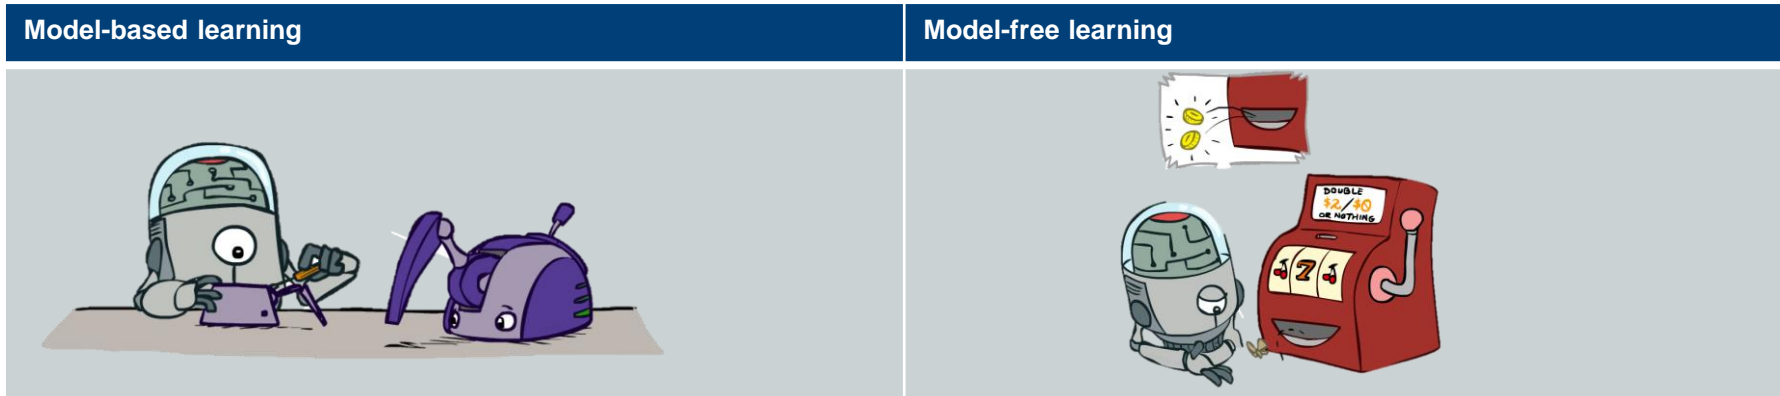
Case Distinctions

Offline Solution (MDPs)	Online Learning (RL)
	
<p>Idea:</p> <ul style="list-style-type: none">• The state space, the action model and the transition model are simulated• Definition as MDP	<p>Idea:</p> <ul style="list-style-type: none">• Agent executes actions (strategy may exist)• Rewards will be experienced• A model may be created
<p>Procedure:</p> <ul style="list-style-type: none">• e.g. Dynamic Programming• Synchronous DP• Asynchronous DP	<p>Procedure:</p> <ul style="list-style-type: none">• Model-based learning• Model-free learning• Active or passive learning

Case Distinctions

Model-based learning	Model-free learning
	
Model-based Idea: <ul style="list-style-type: none">• Learning an approximated model from experience• Solve as if the learned values were correct	Model-free idea: <ul style="list-style-type: none">• Direct learning• Strategy is derived directly• No detour via model identification
Learning of the empirical model: <ul style="list-style-type: none">• Count values s' for all s, a• Normalize to obtain an estimate of the transition model $P_{ss'}^a$• Explore any reward R_s^a	Procedure: <ul style="list-style-type: none">• Active or passive learning

Case Distinctions



Model-based Reinforcement Learning


Model-based RL, Introduction

- **Basic idea:**
 - Learning an approximated transition model based on experience
 - Solving the learned MDP
- **Step 1:** Learning an empirical MDP model
 - Counting the resulting \mathbf{s}' for each state action pair (\mathbf{s}, \mathbf{a})
 - Normalization leads to an estimate of the transitional model $P_{ss'}^a$
 - Observe the reward R_s^a for the episode $(\mathbf{s}, \mathbf{a}, \mathbf{s}')$
- **Step 2:** Solving the learned MDP
 - E.g. Perform value iteration



Model-based RL, Example

In the following GridWorld the transition matrix $P_{ss'}^a$ and the reward function R_s^a are to be learned by RL



Assumption: $\gamma = 1$

Final states:

A, D



Strategy π is given

Learned model parameters

$P_{ss'}^a$		R_s^a	
$\mathbb{P}[C B, East]$	1,00	R_C	-1
$\mathbb{P}[D C, East]$	0,75	R_D	10
$\mathbb{P}[A C, East]$	0,25	R_A	-10
...


Observed episodes (training)

Episode 1				Episode 2				Episode 3				Episode 4			
s	a	s'	R_s^a	s	a	s'	R_s^a	s	a	s'	R_s^a	s	a	s'	R_s^a
B	East	C	-1	B	East	C	-1	E	North	C	-1	E	North	C	-1
C	East	D	-1	C	East	D	-1	C	East	D	-1	C	East	A	-1
D	Exit		+10	D	Exit		+10	D	Exit		+10	A	Exit		-10

Model-free Reinforcement Learning

Model-free Passive Reinforcement Learning


- **Simplified task:** Evaluation of policies
 - **Input:** A fixed policy $\pi(s)$
 - The transition matrix $P_{ss'}^a$ is unknown
 - The reward function R_s^a is unknown
 - **Goal:** Learn the state values $V(s)$



- **In this case:**
 - The learner runs parallel to the policy execution
 - No choice as to which action to perform
 - Just execute the given strategy and learn from experience

Model-free RL, Monte-Carlo Methods

- **Idea:** estimates values based on averaged sample returns
 - Sample Sequences $(s, a, s', R_{ss'})$ by
 - direct interaction with the environment
 - or in simulation
 - Experience is divided into episodes
 - **Episodic Learning:** Update only after completion of an episode
- Learning values from **episodes**
 - **Incremental** – i.e. episode-by-episode (not step-by-step)
 - **Prediction problem:** Calculation of v_π and q_π for fixed, arbitrary strategy π
 - **Strategy improvement**
 - Ideas from Dynamic Programming
- Methods:
 - First-visit MC policy evaluation
 - Every-visit MC policy evaluation



First-visit MC policy evaluation Algorithm

Initialize

$\pi \leftarrow \text{policy to be evaluated}$

$V \leftarrow \text{arbitrary state value function}$

$Returns(s) \leftarrow \text{an empty list}, \forall s \in S$

Repeat (*forever*):

 Generate an episode using π

 For each state s appearing in the episode:

$G \leftarrow \text{return following the first occurrence of } s$

 Add G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$



Model-free RL, Learning from temporal difference (TD)

- **Idea:** Learning from every experience made!
 - Update $V(s)$ directly if a transition $(s, a, s', R_{ss'})$ was executed
 - Since no larger amount of samples is required, the update takes place more frequently
 - **Model-free learning:** No knowledge of MDP transitions / rewards
- Learning values (episodes) from **temporal difference**
 - **Policy** is fixed, **evaluation** will continue
 - Move the values in the direction of the successor state

- Update rule

$$V(s) \leftarrow V(s) + \alpha[R_{ss'} + \gamma V(s') - V(s)]$$

$\alpha := \text{Learning rate parameters}$



$$R_{ss'} + \gamma V(s') := \textbf{TD Goal}$$

$$R_{ss'} + \gamma V(s') - V(s) := \textbf{TD Error}$$

Model-free RL, Learning from temporal difference (TD)

TD-Learning Algorithm

Input: the policy π to be evaluated

Initialize $V(s), \forall s \in S$ *arbitrarily*

Repeat (for each episode e):

 Initialize s

 Repeat (for each step t of the episode):

$a \leftarrow$ given action by π for s


 take action a ; observe reward $R_{ss'}$ and next state s'

$V(s) \leftarrow V(s) + \alpha[R_{ss'} + \gamma V(s') - V(s)]$

$s \leftarrow s'$;

 Until s is final state

Model-free RL, Learning from temporal difference (TD), Example



?



?

$P_{ss'}$

Initial state A_1

$\pi = S \rightarrow A$

$\gamma = 0,9$

$\alpha = 0,5$

$R_{ss'} = ?$

Action executed a		Observed state s'		Observed reward $R_{ss'}$	

$V(s) \leftarrow V(s) + \alpha[R_{ss'} + \gamma V(s') - V(s)]$					$V(s)$	

$V(A_1)$	$V(A_2)$	$V(A_3)$	$V(A_4)$	$V(B_1)$	$V(B_3)$	$V(B_4)$	$V(C_1)$	$V(C_2)$	$V(C_3)$	$V(C_4)$

Model-free RL, Learning from temporal difference (TD), Example

C

B

A


1

2

3

4

?



?

$P_{ss'}$

Initial state A_1

$\pi = S \rightarrow A$

$\gamma = 0,9$

$\alpha = 0,5$

$R_{ss'} = ?$

1. Initialize $V(s), \forall s \in \mathcal{S}$, arbitrarily

Action executed a	Observed state s'	Observed reward $R_{ss'}$

$V(s) \leftarrow V(s) + \alpha[R_{ss'} + \gamma V(s') - V(s)]$	$V(s)$

$V(A_1)$	$V(A_2)$	$V(A_3)$	$V(A_4)$	$V(B_1)$	$V(B_3)$	$V(B_4)$	$V(C_1)$	$V(C_2)$	$V(C_3)$	$V(C_4)$
0	0	0	0	0	0	0	0	0	0	0

Model-free RL, Learning from temporal difference (TD), Example

C

→

→

→

B

↑

↑

A

↑

→

↑

←

1

2

3

4

?

↑

←

?

→

?

$P_{ss'}$

Initial state A_1

$\pi = S \rightarrow A$
 $\gamma = 0,9$
 $\alpha = 0,5$
 $R_{ss'} = ?$

2. Set π for the strategy to be evaluated

Action executed a	Observed state s'	Observed reward $R_{ss'}$

$V(s) \leftarrow V(s) + \alpha[R_{ss'} + \gamma V(s') - V(s)]$	$V(s)$

$V(A_1)$	$V(A_2)$	$V(A_3)$	$V(A_4)$	$V(B_1)$	$V(B_3)$	$V(B_4)$	$V(C_1)$	$V(C_2)$	$V(C_3)$	$V(C_4)$
0	0	0	0	0	0	0	0	0	0	0

Model-free RL, Learning from temporal difference (TD), Example

C

→

→

→

B

↑

↑

A

↑

→

↑

←

1

2

3

4

?

↑

← ?

→ ?

$P_{ss'}$

Initial state A_1

$\pi = S \rightarrow A$

$\gamma = 0,9$

$\alpha = 0,5$

$R_{ss'} = ?$

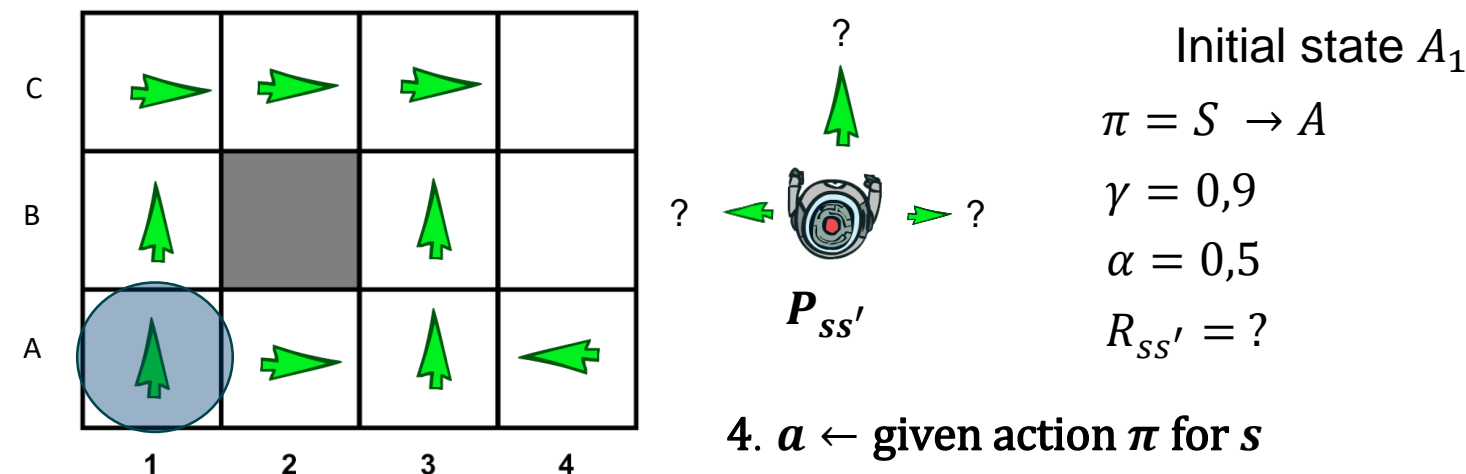
3. Initialize s

Action executed a	Observed state s'						Observed reward $R_{ss'}$			

$V(s) \leftarrow V(s) + \alpha[R_{ss'} + \gamma V(s') - V(s)]$	$V(s)$

$V(A_1)$	$V(A_2)$	$V(A_3)$	$V(A_4)$	$V(B_1)$	$V(B_3)$	$V(B_4)$	$V(C_1)$	$V(C_2)$	$V(C_3)$	$V(C_4)$
0	0	0	0	0	0	0	0	0	0	0

Model-free RL, Learning from temporal difference (TD), Example



Action executed a	Observed state s'	Observed reward $R_{ss'}$
a_{South}		

$V(s) \leftarrow V(s) + \alpha[R_{ss'} + \gamma V(s') - V(s)]$	$V(s)$

$V(A_1)$	$V(A_2)$	$V(A_3)$	$V(A_4)$	$V(B_1)$	$V(B_3)$	$V(B_4)$	$V(C_1)$	$V(C_2)$	$V(C_3)$	$V(C_4)$
0	0	0	0	0	0	0	0	0	0	0

Model-free RL, Learning from temporal difference (TD), Example

C

B

A

→

→

→

↑

↑

↑

→

↑

←

1

2

3


4

?

↑

?

←



→

?

$P_{ss'}$

Initial state A_1

$\pi = S \rightarrow A$

$\gamma = 0,9$

$\alpha = 0,5$

$R_{ss'} = ?$

5. Execute a ; observe reward $R_{ss'}$
and next state s'

Action executed a	Observed state s'	Observed reward $R_{ss'}$
a_{North}	B_1	-0,05

$V(s) \leftarrow V(s) + \alpha[R_{ss'} + \gamma V(s') - V(s)]$	$V(s)$

$V(A_1)$	$V(A_2)$	$V(A_3)$	$V(A_4)$	$V(B_1)$	$V(B_3)$	$V(B_4)$	$V(C_1)$	$V(C_2)$	$V(C_3)$	$V(C_4)$
0	0	0	0	0	0	0	0	0	0	0

Model-free RL, Learning from temporal difference (TD), Example

C

B

A

→

→

→

↑

↑

↑

→

↑

←

1

2

3


4

?

↑

?

←



→

?

$P_{ss'}$

Initial state A_1

$\pi = S \rightarrow A$

$\gamma = 0,9$

$\alpha = 0,5$

$R_{ss'} = ?$


6. $V(s) \leftarrow V(s) + \alpha[R_{ss'} + \gamma V(s') - V(s)]$

Action executed a	Observed state s'	Observed reward $R_{ss'}$
a_{North}	B_1	$-0,05$

$V(s) \leftarrow V(s) + \alpha[R_{ss'} + \gamma V(s') - V(s)]$	$V(s)$
$0 + 0,5[-0,05 + 0,9 * 0 - 0]$	$-0,025$

$V(A_1)$	$V(A_2)$	$V(A_3)$	$V(A_4)$	$V(B_1)$	$V(B_3)$	$V(B_4)$	$V(C_1)$	$V(C_2)$	$V(C_3)$	$V(C_4)$
$-0,025$	0	0	0	0	0	0	0	0	0	0

Model-free RL, Learning from temporal difference (TD), Example




Action executed a	Observed state s'	Observed reward $R_{ss'}$
a_{North}	B_1	$-0,05$


$V(s) \leftarrow V(s) + \alpha[R_{ss'} + \gamma V(s') - V(s)]$	$V(s)$
$0 + 0,5[-0,05 + 0,9 * 0 - 0]$	$-0,025$

$V(A_1)$	$V(A_2)$	$V(A_3)$	$V(A_4)$	$V(B_1)$	$V(B_3)$	$V(B_4)$	$V(C_1)$	$V(C_2)$	$V(C_3)$	$V(C_4)$
$-0,025$	0	0	0	0	0	0	0	0	0	0

Model-free RL, Learning from temporal difference (TD), Example




Action executed a				Observed state s'			Observed reward $R_{ss'}$			
a_{North}										
$V(s) \leftarrow V(s) + \alpha[R_{ss'} + \gamma V(s') - V(s)]$							$V(s)$			
$V(A_1)$	$V(A_2)$	$V(A_3)$	$V(A_4)$	$V(B_1)$	$V(B_3)$	$V(B_4)$	$V(C_1)$	$V(C_2)$	$V(C_3)$	$V(C_4)$
-0,025	0	0	0	0	0	0	0	0	0	0


$$R_{ss'} = ?$$

Action executed \boldsymbol{a}	Observed state \boldsymbol{s}'	Observed reward $R_{ss'}$
\boldsymbol{a}_{North}	B_1	$-0,05$


[illegible]


$$R_{ss'} = ?$$

Action executed a	Observed state s'	Observed reward $R_{ss'}$
a_{North}	B_1	$-0,05$

$V(A_1)$	$V(A_2)$	$V(A_3)$	$V(A_4)$	$V(B_1)$	$V(B_3)$	$V(B_4)$	$V(C_1)$	$V(C_2)$	$V(C_3)$	$V(C_4)$
-0,025	0	0	0	-0,025	0	0	0	0	0	0

Model-free RL, Learning from temporal difference (TD), Example




Action executed a	Observed state s'	Observed reward $R_{ss'}$
a_{North}	B_1	$-0,05$

$V(s) \leftarrow V(s) + \alpha[R_{ss'} + \gamma V(s') - V(s)]$	$V(s)$
$0 + 0,5[-0,05 + 0,9 * 0 - 0]$	$-0,025$

$V(A_1)$	$V(A_2)$	$V(A_3)$	$V(A_4)$	$V(B_1)$	$V(B_3)$	$V(B_4)$	$V(C_1)$	$V(C_2)$	$V(C_3)$	$V(C_4)$
$-0,025$	0	0	0	$-0,025$	0	0	0	0	0	0

Model-free Active Reinforcement Learning




Model-free (Active) RL, Q-Learning

- **Idea:** Learning an action/value representation
 - During TD learning a value representation was learned
 - Neither learning nor action selection requires a model $P_{ss'}^a$,
 - Therefore, this is a model-free method
- **Procedure**
 - No explicit policy (unlike TD learning) \Rightarrow active
 - Algorithm always selects the **action with best Q-value**
- **Update-Rule**

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[R_{ss'} + \gamma \max_a Q(s', a) - Q(s, a) \right]$$

$\alpha :=$ learning rate parameter



Model-free (Active) RL, Q-Learning

Q-Learning Algorithm

Initialize $Q(s, a), \forall s \in S, a \in A(s)$ arbitrarily, and $Q(\text{Final State}, \cdot) = 0$

Repeat (for each episode e):

Initialize s

Repeat (for each step t of the episode):

Choose a from s from the policy derived from Q

Take action a and observe $R_{ss'}, s'$


$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[R_{ss'} + \gamma \max_a Q(s', a) - Q(s, a) \right]$$

$s \leftarrow s'$;


Until s is final state

Model-free (Active) RL, Q-Learning

- Impressive results :
 - Q-Learning converges to an optimal strategy - even if the agent behaves suboptimally
- This is called **Off-Policy Learning**
- Conditional:
 - There has been enough exploration
 - The learning rate parameter may have to decrease over time.
 - ...but this must not happen too quickly.
 - In a nutshell: After enough time, it is not relevant how actions are selected for exploration (!)




Model-free (Active) RL, Exploration vs. Exploitation




Model-free (Active) RL, Exploration vs. Exploitation

- **Easiest: Random actions (ϵ -greedy)**
 - In every time step, toss a coin
 - With (little) probability ϵ : choose a random action
 - With (high) probability $1 - \epsilon$: act on the basis of the current strategy
- **Problem with ϵ -greedy**
 - Even though the entire state / action space is explored, the agent may still behave suboptimally afterwards
 - A Solution: decrease ϵ over time



Further Reading Material

- Berkeley Course – Artificial Intelligence (<https://inst.eecs.berkeley.edu/~cs188/fa11/lectures.html>)
- David Silver – Reinforcement Learning (<http://www0.cs.ucl.ac.uk/staff/d.silver/web/Teaching.html>)



Thank you for your attention!

Lecture Team AIDAE
aidae@ima-ifu.rwth-aachen.de